

# Chenggang Wu

465 Soda Hall, Berkeley, CA 94720, USA  
cgwu@berkeley.edu • cgwu.io

## EDUCATION

### University of California, Berkeley

- Doctor of Philosophy (Ph.D.) in Computer Science Aug 2015 – Present
  - Adviser: Joseph M. Hellerstein
  - Research areas: Data-centric Systems, Distributed Systems, Machine Learning

### Brown University

- Bachelor of Science (B.S.) in Computer Science May 2015
  - Magna cum laude with Honors in Computer Science.
  - Cumulative GPA: 4.0 / 4.0

## RESEARCH EXPERIENCE

### RISE Lab, UC Berkeley

Sep 2016 – Present

- Fluent: a declarative programming framework for building distributed systems.
- Anna: a super-fast, flexibly consistent, autoscaling and cost-effective Key-Value Store.

### CISL Lab, Microsoft

May 2017 – Oct 2018

- CardLearner: a system that learns cardinality models from previous query executions and uses them to aid the query optimizer in predicting the cardinalities of future queries.

### Data Management Research Group, Brown University

Jan 2014 – May 2015

- S-Store: a transactional stream processing system that simultaneously accommodates OLTP workload and streaming applications.
- Seer: a predictive prefetching and caching middleware for supporting exploratory visualization of big data.

## PUBLICATIONS

- [1] **Chenggang Wu**, Jose Faleiro, Yihan Lin, Joseph M. Hellerstein, “Anna: A KVS For Any Scale,” *TKDE 2019 (to appear)*, “[Best of ICDE 2018](#)” Special Issue.
- [2] **Chenggang Wu**, Vikram Sreekanti, Joseph M. Hellerstein, “Eliminating Boundaries in Cloud Storage with Anna,” *VLDB 2019 (to appear)*.
- [3] Joseph M. Hellerstein, Jose Faleiro, Joseph E. Gonzalez, Johann Schleier-Smith, Vikram Sreekanti, Alexey Tumanov, **Chenggang Wu**, “Serverless Computing: One Step Forward, Two Steps Back,” *CIDR 2019*.
- [4] **Chenggang Wu**, Alekh Jindal, Saeed Amizadeh, Hiren Patel, Wangchao Le, Shi Qiao, Sriram Rao, “Towards a Learning Optimizer for Shared Clouds,” *VLDB 2019*.
- [5] **Chenggang Wu**, Jose Faleiro, Yihan Lin, Joseph M. Hellerstein, “Anna: A KVS For Any Scale,” *ICDE 2018*, [Best of Conference Selection](#).

## INVITED TALKS

### “Eliminating Boundaries in Cloud Storage with Anna”

- Alibaba Group, October 2018
- RISE Lab Retreat, May 2018
- Amazon Web Services, May 2018

### “Towards a Learning Optimizer for Shared Clouds”

- Microsoft, August 2017

“Anna: A KVS For Any Scale”

- Alibaba Group, October 2018
- RISE Lab Retreat, May 2018
- Amazon Web Services, May 2018
- ICDE 2018, April 2018

**PROFESSIONAL  
SERVICES**

Reviewer for TKDE 2019.

U.C. Berkeley EECS Graduate Admissions Committee Member.

Organizer of Database Seminar for Berkeley Database Group.

**TEACHING**

**Graduate Student Instructor**, UC Berkeley

- CS186/286A: Introduction to Database Systems, Spring 2018
  - with Prof. Mehul Shah
- CS186/286A: Introduction to Database Systems, Spring 2016
  - with Prof. Joseph E. Gonzalez and Prof. Joseph M. Hellerstein

**Head Teaching Assistant**, Brown University

- CS1951A: Introduction to Data Science, Spring 2015
  - with Prof. Tim Kraska
- CS1270: Database Management Systems, Fall 2014
  - with Prof. Stan Zdonik

**ACADEMIC  
AWARDS**

Computing Research Association Outstanding Undergraduate Researchers Award - Nominee,  
Brown University

- Awarded to the top three students in the computer science department based on research contributions.

**TECHNICAL  
BACKGROUND**

Advanced Coursework: Artificial Intelligence, Databases, Data Science, Distributed Systems, Machine Learning, Parallel Computing, Programming Languages & Compilers, Algorithms , Systems Security.

Programming: C/C++, Python, Java, JavaScript, Scala.